

CODE:

```
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from sklearn.model_selection import LeaveOneGroupOut
from transformers import BertTokenizer, BertModel
from sklearn.metrics import accuracy_score, f1_score
import os
import re
import cv2 # OpenCV for video processing
import face_recognition # Now mandatory
import librosa # For audio analysis
from moviepy.editor import VideoFileClip # For extracting audio
import torchvision.models as models
import torchvision.transforms as transforms
from PIL import Image
import time
import warnings

# Suppress warnings from l
# libraries like moviepy/librosa if needed
warnings.filterwarnings("ignore")

# --- Configuration ---
BERT_MODEL_NAME = 'bert-base-uncased'
VISUAL_MODEL_NAME = 'resnet18' # Using ResNet18 for visual features
AUDIO_N_MFCC = 13 # Number of MFCCs for audio features
VISUAL_FRAMES_TO_SAMPLE = 70 # Number of frames to sample per video for visual
features
DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {DEVICE}")

# --- Data Directory Setup ---
BASE_DIR = os.getcwd() # Get current working directory
data_dir = os.path.join(BASE_DIR, 'Real-life Deception Detection 2016')
annotation_file = os.path.join(data_dir, 'Annotation', 'All_Gestures_Deceptive and
Truthful.csv')
# --- 1. Subject Identification (Mandatory Facial Recognition) ---
def identify_subjects_facial_recognition(data):
    """
    Identifies subjects MANDATORILY using facial recognition from the first frame.
    Maps trial_id to a subject label.

    Args:
        data (list): List of dictionaries from load_data, containing 'video_path'
and 'video_id' (trial_id).

    Returns:
        dict: A dictionary mapping trial_ids (video_ids) to subject labels (e.g.,
'subject_1', 'unknown_trial_xyz').
    """
    print("Starting mandatory facial recognition for subject identification...")
    subject_mapping = {}
    known_faces = {} # Store known face encodings and labels {label: encoding}
    subject_counter = 1
    unknown_counter = 1
```

```

for item in data:
    video_path = item['video_path']
    trial_id = item['video_id'] # Use the actual trial_id
    subject_label = None

    try:
        cap = cv2.VideoCapture(video_path)
        if not cap.isOpened():
            print(f"Warning: Could not open video file {video_path} for trial
{trial_id}. Assigning unknown subject.")
            subject_label = f'unknown_video_open_error_{unknown_counter}'
            unknown_counter += 1
            subject_mapping[trial_id] = subject_label
            continue

        # Read the first frame
        ret, frame = cap.read()
        if not ret:
            print(f"Warning: Could not read frame from video {video_path} for
trial {trial_id}. Assigning unknown subject.")
            subject_label = f'unknown_frame_read_error_{unknown_counter}'
            unknown_counter += 1
            subject_mapping[trial_id] = subject_label
            cap.release()
            continue

        # Convert the frame to RGB (face_recognition uses RGB)
        rgb_frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)

        # Find face locations and encodings in the frame
        face_locations = face_recognition.face_locations(rgb_frame)
        face_encodings = face_recognition.face_encodings(rgb_frame,
face_locations)

        if not face_encodings:
            # print(f"Warning: No faces found in the first frame of video
{trial_id}. Assigning unique unknown subject.")
            subject_label = f'unknown_no_face_{unknown_counter}' # Assign a
unique unknown label based on trial id
            unknown_counter += 1
        else:
            # Use the first face found
            current_face_encoding = face_encodings[0]

            # Check for matches with known faces
            match_found = False
            known_labels = list(known_faces.keys())
            if known_labels:
                known_encodings = list(known_faces.values())
                # Increase tolerance slightly if needed, default is 0.6
                matches = face_recognition.compare_faces(known_encodings,
current_face_encoding, tolerance=0.6)
                # Find the first match
                try:
                    first_match_index = matches.index(True)
                    subject_label = known_labels[first_match_index]
                    match_found = True
                except ValueError: # No True value found in matches
                    pass

```

```

        if not match_found:
            # If no match, add the face to known faces
            subject_label = f'subject_{subject_counter}'
            known_faces[subject_label] = current_face_encoding
            subject_counter += 1

        subject_mapping[trial_id] = subject_label
        cap.release()

    except Exception as e:
        print(f"Error processing video {video_path} for trial {trial_id}: {e}.
Assigning unknown subject.")
        subject_label = f'unknown_processing_error_{unknown_counter}'
        unknown_counter += 1
        subject_mapping[trial_id] = subject_label
        if 'cap' in locals() and cap.isOpened():
            cap.release()

    print(f"Facial recognition complete. Identified {subject_counter - 1} unique
subjects and {unknown_counter - 1} videos needing unique IDs.")
    return subject_mapping

# --- 2. Data Loading (Slightly modified for clarity) ---
def load_data(data_dir, annotation_file):
    """Loads and synchronizes annotation, transcription, and video data."""
    print("Loading data...")

    # Construct paths
    clip_dirs = [
        os.path.join(data_dir, 'Clips', 'Deceptive'),
        os.path.join(data_dir, 'Clips', 'Truthful')
    ]
    transcript_dirs = [
        os.path.join(data_dir, 'Transcription', 'Deceptive'),
        os.path.join(data_dir, 'Transcription', 'Truthful')
    ]

    # Initialize lists to store paths
    video_paths = []
    transcription_paths = []

    # Load video paths
    for clip_dir in clip_dirs:
        if os.path.isdir(clip_dir):
            for filename in os.listdir(clip_dir):
                if filename.endswith(".mp4"):
                    video_paths.append(os.path.join(clip_dir, filename))
        else:
            print(f"Warning: Clip directory not found: {clip_dir}")

    # Load transcription paths - UPDATED to look for .txt files
    for transcript_dir in transcript_dirs:
        if os.path.isdir(transcript_dir):
            for filename in os.listdir(transcript_dir):
                if filename.endswith(".txt"): # Changed from .csv to .txt
                    transcription_paths.append(os.path.join(transcript_dir,
filename))

```

```

        else:
            print(f"Warning: Transcription directory not found: {transcript_dir}")

# Load annotations
try:
    annotations_df = pd.read_csv(annotation_file)
except FileNotFoundError:
    print(f"Error: Annotation file not found at {annotation_file}")
    return []

# Create mapping dictionaries
video_path_dict = {}
for video_path in video_paths:
    video_filename = os.path.basename(video_path)
    match = re.search(r"trial_(truth|lie)_(\d+)\.mp4", video_filename)
    if match:
        trial_id = f"trial_{match.group(1)}_{match.group(2)}"
        video_path_dict[trial_id] = video_path

# Print debug information
print(f"Found {len(video_paths)} video files")
print(f"Found {len(transcription_paths)} transcription files")

transcription_dict = {}
for transcript_path in transcription_paths:
    try:
        # Read the .txt file directly
        with open(transcript_path, 'r', encoding='utf-8') as f:
            transcription_text = f.read().strip()

        # Extract trial_id from filename
        filename = os.path.basename(transcript_path)
        trial_id = filename.replace('.txt', '') # Remove .txt extension

        print(f"Reading transcription file: {transcript_path}")
        transcription_dict[trial_id] = transcription_text
        print(f"Added transcription for {trial_id}")

    except Exception as e:
        print(f"Error reading transcription file {transcript_path}: {e}")

# Synchronize data
synchronized_data = []
processed_ids = set()

if 'id' not in annotations_df.columns or 'class' not in
annotations_df.columns:
    print(f"Error: Annotation file missing 'id' or 'class' column.")
    return []

for _, row in annotations_df.iterrows():
    trial_id = str(row['id']).replace('.mp4', '')
    annotation_label = row['class']

    if annotation_label not in ['truthful', 'deceptive']:
        print(f"Warning: Skipping trial {trial_id} due to unexpected class
label: {annotation_label}")
        continue

```

```

        if trial_id in processed_ids:
            print(f"Warning: Duplicate trial ID {trial_id} found. Skipping.")
            continue

        transcription_text = transcription_dict.get(trial_id)
        video_path = video_path_dict.get(trial_id)

        if transcription_text is None:
            print(f"Warning: Transcription not found for trial {trial_id}")
            continue
        if video_path is None:
            print(f"Warning: Video path not found for trial {trial_id}")
            continue

        # Map labels to numerical values
        label_map = {'truthful': 0, 'deceptive': 1}
        numeric_label = label_map.get(annotation_label)
        if numeric_label is None:
            print(f"Warning: Could not map label '{annotation_label}' for trial
{trial_id}")
            continue

        synchronized_data.append({
            'annotation': numeric_label,
            'transcription': transcription_text,
            'video_id': trial_id,
            'video_path': video_path
        })
        processed_ids.add(trial_id)

    print(f"Data loading complete. Found {len(synchronized_data)} synchronized
    trials.")
    return synchronized_data

# --- 3. Feature Extraction ---

# 3.1 NLP Feature Extraction (BERT) - Unchanged conceptually
def extract_nlp_features(transcriptions):
    """ Extracts BERT embeddings for a list of transcriptions. """
    print("Extracting NLP features (BERT)...")
    tokenizer = BertTokenizer.from_pretrained(BERT_MODEL_NAME)
    model = BertModel.from_pretrained(BERT_MODEL_NAME).to(DEVICE)
    model.eval()
    nlp_features = []
    with torch.no_grad():
        for i, text in enumerate(transcriptions):
            try:
                # Ensure text is a string
                text = str(text) if text is not None else ""
                if not text.strip(): # Handle empty strings
                    print(f"Warning: Empty transcription for item {i}. Using zero
vector.")

                # Get expected hidden size from model config
                hidden_size = model.config.hidden_size
                sentence_embedding = np.zeros((1, hidden_size))
            else:
                inputs = tokenizer(text, return_tensors='pt',

```

```

                                truncation=True, padding=True,
max_length=512).to(DEVICE) # Added max_length
                                outputs = model(**inputs)
                                # Mean of the last hidden state
                                sentence_embedding =
outputs.last_hidden_state.mean(dim=1).cpu().numpy() # (1, hidden_size)
                                nlp_features.append(sentence_embedding)
                                except Exception as e:
                                    print(f"Error extracting NLP features for item {i}: {e}. Using
zero vector.")
                                    hidden_size = model.config.hidden_size
                                    nlp_features.append(np.zeros((1, hidden_size)))

                                print("NLP feature extraction complete.")
                                # Ensure all features are arrays and handle potential shape issues before
stacking
                                processed_features = []
                                target_shape = None
                                for feat in nlp_features:
                                    if isinstance(feat, np.ndarray):
                                        if target_shape is None:
                                            target_shape = feat.shape
                                        # If shape mismatch, pad or truncate (or use zeros as done in
exception handling)
                                        if feat.shape != target_shape:
                                            print(f"Warning: NLP feature shape mismatch {feat.shape} vs
{target_shape}. Using zero vector.")
                                            processed_features.append(np.zeros(target_shape))
                                        else:
                                            processed_features.append(feat)
                                    else: # Should not happen if exceptions are caught, but as safeguard
                                        if target_shape is None: # Need a shape defined first
                                            raise ValueError("Cannot process non-array NLP feature without a
target shape.")
                                        print(f"Warning: Non-array NLP feature found. Using zero vector.")
                                        processed_features.append(np.zeros(target_shape))

                                if not processed_features:
                                    return np.array([]) # Return empty array if no features were processed

                                return np.vstack(processed_features) # (num_trials, hidden_size)

# 3.2 Audio Feature Extraction (MFCCs)
def extract_audio_features(video_paths):
    """ Extracts MFCC features from the audio track of video files. """
    print("Extracting Audio features (MFCCs)...")
    audio_features = []
    temp_audio_dir = "temp_audio"
    if not os.path.exists(temp_audio_dir):
        os.makedirs(temp_audio_dir)

    num_features = AUDIO_N_MFCC * 2 # Mean and Std Dev for each MFCC

    for i, video_path in enumerate(video_paths):
        start_time = time.time()
        temp_audio_path = os.path.join(temp_audio_dir, f"temp_{i}.wav")

```

```

        feature_vector = np.zeros(num_features) # Default to zeros

    try:
        # Extract audio using moviepy
        with VideoFileClip(video_path) as video_clip:
            if video_clip.audio is None:
                print(f"Warning: Video {i} ({os.path.basename(video_path)})
has no audio track. Using zeros.")
            else:
                video_clip.audio.write_audiofile(temp_audio_path,
codec='pcm_s16le', logger=None) # Use logger=None to reduce console output

                # Load audio and extract MFCCs using librosa
                y, sr = librosa.load(temp_audio_path, sr=None) # Load with native
sample rate
                if len(y) > 0: # Check if audio signal is not empty
                    mfccs = librosa.feature.mfcc(y=y, sr=sr, n_mfcc=AUDIO_N_MFCC)
                    mfccs_mean = np.mean(mfccs, axis=1)
                    mfccs_std = np.std(mfccs, axis=1)
                    feature_vector = np.concatenate((mfccs_mean, mfccs_std))
                else:
                    print(f"Warning: Audio signal empty after extraction for video
{i}. Using zeros.")

    except Exception as e:
        print(f"Error extracting audio features for video {i}
({os.path.basename(video_path)}): {e}. Using zeros.")
        # Ensure feature_vector remains zeros

    finally:
        # Clean up temporary audio file
        if os.path.exists(temp_audio_path):
            try:
                os.remove(temp_audio_path)
            except Exception as e:
                print(f"Warning: Could not remove temp audio file
{temp_audio_path}: {e}")

        audio_features.append(feature_vector)
        # print(f"Processed audio for video {i+1}/{len(video_paths)} in
{time.time() - start_time:.2f}s") # Optional progress

    # Clean up temp directory if empty
    try:
        if not os.listdir(temp_audio_dir):
            os.rmdir(temp_audio_dir)
    except Exception as e:
        print(f"Warning: Could not remove temp audio directory {temp_audio_dir}:
{e}")

print("Audio feature extraction complete.")
return np.array(audio_features) # (num_trials, num_audio_features)

# 3.3 Visual Feature Extraction (ResNet)
def extract_visual_features(video_paths):

```

```

""" Extracts aggregated visual features using a pre-trained ResNet model. """
print("Extracting Visual features (ResNet)...")

# Load pre-trained ResNet model without the final classification layer
vis_model = models.resnet18(pretrained=True)
vis_model = nn.Sequential(*list(vis_model.children())[:-1]) # Remove the fully
connected layer
vis_model = vis_model.to(DEVICE)
vis_model.eval()

# Define image transformations appropriate for ResNet
preprocess = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
0.225]),
])

visual_features = []
num_visual_features = 512 # ResNet18 output size before FC layer

with torch.no_grad():
    for i, video_path in enumerate(video_paths):
        start_time = time.time()
        video_feature_vector = np.zeros(num_visual_features) # Default to
zeros
        try:
            cap = cv2.VideoCapture(video_path)
            if not cap.isOpened():
                print(f"Warning: Could not open video {i}
({os.path.basename(video_path)}). Using zeros.")
                visual_features.append(video_feature_vector)
                continue

            frame_count = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
            if frame_count <= 0:
                print(f"Warning: Video {i} has no frames or invalid frame
count. Using zeros.")
                visual_features.append(video_feature_vector)
                cap.release()
                continue

            frame_indices = np.linspace(0, frame_count - 1,
VISUAL_FRAMES_TO_SAMPLE, dtype=int) # Sample frames evenly

            frames_data = []
            for frame_index in frame_indices:
                cap.set(cv2.CAP_PROP_POS_FRAMES, frame_index)
                ret, frame = cap.read()
                if ret:
                    # Convert frame BGR -> RGB -> PIL Image -> Apply
transforms
                    frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
                    img_pil = Image.fromarray(frame_rgb)
                    img_tensor = preprocess(img_pil).unsqueeze(0).to(DEVICE) #
Add batch dimension
                    frames_data.append(img_tensor)

```



```

        # else: # Optional: Warn if specific frame fails
        #     print(f"Warning: Could not read frame {frame_index} from
video {i}.")

    cap.release()

    if frames_data:
        # Stack frame tensors and pass through the model
        batch_tensor = torch.cat(frames_data, dim=0)
        frame_outputs = vis_model(batch_tensor) # (num_sampled_frames,
num_visual_features, 1, 1)
        frame_outputs = frame_outputs.squeeze() # Remove trailing 1s -
> (num_sampled_frames, num_visual_features)
        # Aggregate features (e.g., mean pooling)
        video_feature_vector = torch.mean(frame_outputs,
dim=0).cpu().numpy()
    else:
        print(f"Warning: No frames could be processed for video {i}.
Using zeros.")
        # video_feature_vector remains zeros

except Exception as e:
    print(f"Error extracting visual features for video {i}
({os.path.basename(video_path)}): {e}. Using zeros.")
    if 'cap' in locals() and cap.isOpened():
        cap.release()
    # video_feature_vector remains zeros

    visual_features.append(video_feature_vector)
    # print(f"Processed visual for video {i+1}/{len(video_paths)} in
{time.time() - start_time:.2f}s")

print("Visual feature extraction complete.")
return np.array(visual_features) # (num_trials, num_visual_features)

# --- 4. Data Preparation for LOSO ---
def prepare_loso(data, subject_mapping):
    """
    Prepares data for LOSO cross-validation using the mandatory subject mapping.
    Returns raw transcriptions for later NLP processing.
    """
    print("Preparing data for LOSO...")
    annotations = [item['annotation'] for item in data]
    transcriptions = [item['transcription'] for item in data] # Keep raw text
    video_ids = [item['video_id'] for item in data]

    # Map video_ids (trial_ids) to subject IDs using the facial recognition
mapping
    mapped_subject_ids = []
    valid_indices = [] # Keep track of trials with successful subject mapping
    for idx, video_id in enumerate(video_ids):
        subject_id = subject_mapping.get(video_id)
        if subject_id is None:
            print(f"Critical Warning: No subject mapping found for video_id
{video_id}. This trial will be skipped in prepare_loso.")

```

```

        else:
            mapped_subject_ids.append(subject_id)
            valid_indices.append(idx)

    if len(valid_indices) < len(data):
        print(f"Warning: {len(data) - len(valid_indices)} trials were skipped due
to missing subject mapping.")

    # Filter data based on valid indices
    annotations = np.array(annotations)[valid_indices]
    transcriptions = [transcriptions[i] for i in valid_indices]
    mapped_subject_ids = [mapped_subject_ids[i] for i, _ in
enumerate(valid_indices)] # Already filtered conceptually

    print(f"Data preparation complete. {len(annotations)} trials ready for LOSO.")
    # Note: Audio/Visual features are extracted separately AFTER prepare_loso
filters trials
    return annotations, transcriptions, mapped_subject_ids, valid_indices

# --- 5. Multimodal Model Implementation ---
class MultimodalDeceptionModel(nn.Module):
    """Enhanced multimodal model with sophisticated HSTA and NLP processing."""
    def __init__(self, nlp_input_size, audio_input_size, visual_input_size,
hidden_size, num_classes):
        super(MultimodalDeceptionModel, self).__init__()

        # NLP processor (unchanged)
        self.nlp_processor = nn.Sequential(
            nn.Linear(nlp_input_size, hidden_size),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(hidden_size, hidden_size),
            nn.ReLU()
        )

        # Audio processor (unchanged)
        self.audio_processor = nn.Sequential(
            nn.Linear(audio_input_size, hidden_size),
            nn.ReLU(),
            nn.Dropout(0.3),
            nn.Linear(hidden_size, hidden_size),
            nn.ReLU()
        )

        # Enhanced visual processor with HSTA
        self.visual_processor = nn.Sequential(
            nn.Linear(visual_input_size, hidden_size),
            nn.ReLU(),
            nn.Dropout(0.3)
        )
        self.hsta = HierarchicalSpatioTemporalAttention(
            input_size=hidden_size,
            hidden_size=hidden_size
        )

        # Feature fusion with attention

```

```

        self.fusion_attention = nn.MultiheadAttention(
            embed_dim=hidden_size * 3,
            num_heads=8,
            dropout=0.1
        )

        # Final classifier
        self.classifier = nn.Sequential(
            nn.Linear(hidden_size * 3, hidden_size),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(hidden_size, num_classes)
        )

    def forward(self, nlp_data, audio_data, visual_data):
        # Process each modality
        nlp_processed = self.nlp_processor(nlp_data)
        audio_processed = self.audio_processor(audio_data)

        # Process visual data with HSTA
        visual_processed = self.visual_processor(visual_data)
        # Reshape for HSTA (batch_size, num_frames, hidden_size)
        batch_size = visual_processed.size(0)
        visual_processed = visual_processed.view(batch_size, -1,
visual_processed.size(-1))
        visual_processed = self.hsta(visual_processed)
        # Take the last frame's representation
        visual_processed = visual_processed[:, -1, :]

        # Concatenate features
        fused_features = torch.cat((nlp_processed, audio_processed,
visual_processed), dim=1)

        # Apply fusion attention
        fused_features = fused_features.unsqueeze(0) # Add sequence dimension
        fused_features, _ = self.fusion_attention(fused_features, fused_features,
fused_features)
        fused_features = fused_features.squeeze(0)

        # Final classification
        output = self.classifier(fused_features)
        return output

# --- 6. Training and Evaluation (Modified for Multimodal) ---
def train_evaluate(model,
                    nlp_train, audio_train, visual_train, labels_train,
                    nlp_test, audio_test, visual_test, labels_test,
                    optimizer, criterion, device, epoch):
    """ Trains and evaluates the multimodal model for one epoch. """
    model.train()
    optimizer.zero_grad()

    # Move training data to device
    nlp_train_tensor = torch.FloatTensor(nlp_train).to(device)
    audio_train_tensor = torch.FloatTensor(audio_train).to(device)
    visual_train_tensor = torch.FloatTensor(visual_train).to(device)
    labels_train_tensor = torch.LongTensor(labels_train).to(device)

    # Forward pass (Training)

```

```

outputs = model(nlp_train_tensor, audio_train_tensor, visual_train_tensor)
loss = criterion(outputs, labels_train_tensor)

# Backward pass and optimization
loss.backward()
optimizer.step()

# Evaluation
model.eval()
with torch.no_grad():
    # Move test data to device
    nlp_test_tensor = torch.FloatTensor(nlp_test).to(device)
    audio_test_tensor = torch.FloatTensor(audio_test).to(device)
    visual_test_tensor = torch.FloatTensor(visual_test).to(device)
    labels_test_tensor = torch.LongTensor(labels_test).to(device)

    # Forward pass (Evaluation)
    outputs_test = model(nlp_test_tensor, audio_test_tensor,
visual_test_tensor)
    _, predicted = torch.max(outputs_test.data, 1)

    # Calculate metrics
    labels_test_cpu = labels_test_tensor.cpu().numpy()
    predicted_cpu = predicted.cpu().numpy()
    accuracy = accuracy_score(labels_test_cpu, predicted_cpu)
    f1 = f1_score(labels_test_cpu, predicted_cpu, average='weighted',
zero_division=0) # Added zero_division

    return accuracy, f1, loss.item()

# --- 7. Run LOSO Cross-Validation (Modified for Multimodal) ---
def run_loso(annotations, nlp_features, audio_features, visual_features,
subject_ids,
            checkpoint_dir="checkpoints", num_epochs=50, learning_rate=0.001,
hidden_size=128):
    """ Runs LOSO cross-validation for the multimodal model. """
    print("Starting LOSO Cross-Validation...")
    loso = LeaveOneGroupOut()
    all accuracies = []
    all_f1s = []
    all_losses = [] # To store loss per fold and seed
    num_seeds = 3 # Keep number of seeds

    # Create checkpoint directory if it doesn't exist
    if not os.path.exists(checkpoint_dir):
        os.makedirs(checkpoint_dir)

    num_classes = len(np.unique(annotations))
    if num_classes < 2:
        print(f"Error: Only {num_classes} unique class found. Cannot perform
classification.")
        return

    # Use subject_ids (mapped) as groups for LOSO
    # Ensure all feature arrays have the same number of samples as
annotations/subject_ids
    n_samples = len(annotations)

```

```

    assert len(nlp_features) == n_samples, f"NLP features length mismatch:
{len(nlp_features)} vs {n_samples}"
    assert len(audio_features) == n_samples, f"Audio features length mismatch:
{len(audio_features)} vs {n_samples}"
    assert len(visual_features) == n_samples, f"Visual features length mismatch:
{len(visual_features)} vs {n_samples}"
    assert len(subject_ids) == n_samples, f"Subject IDs length mismatch:
{len(subject_ids)} vs {n_samples}"

    fold_num = 0
    for train_index, test_index in loso.split(nlp_features, annotations,
groups=subject_ids):
        fold_num += 1
        fold_accuracies_seeds = []
        fold_f1s_seeds = []
        fold_losses_seeds = {} # Store losses per seed {seed: [epoch_losses]}

        test_subject = np.unique(np.array(subject_ids)[test_index])[0]
        print(f"\n--- Fold {fold_num}/{loso.get_n_splits(groups=subject_ids)}:
Testing on Subject {test_subject} ---")

        # Split data for this fold
        nlp_train, nlp_test = nlp_features[train_index], nlp_features[test_index]
        audio_train, audio_test = audio_features[train_index],
audio_features[test_index]
        visual_train, visual_test = visual_features[train_index],
visual_features[test_index]
        labels_train, labels_test = annotations[train_index],
annotations[test_index]

        # Check if train or test set is empty for this fold (can happen with LOSO
if subject has few samples)
        if len(labels_train) == 0 or len(labels_test) == 0:
            print(f"Warning: Skipping Fold {fold_num} due to empty train
({len(labels_train)}) or test ({len(labels_test)}) set.")
            continue

        # Get feature dimensions dynamically
        nlp_dim = nlp_train.shape[1]
        audio_dim = audio_train.shape[1]
        visual_dim = visual_train.shape[1]

        for seed in range(num_seeds):
            print(f" Seed {seed + 1}/{num_seeds}")
            torch.manual_seed(seed)
            np.random.seed(seed) # Also seed numpy if any random operations happen
there

        # Model initialization, optimizer, and loss function
        model = MultimodalDeceptionModel(
            nlp_input_size=nlp_dim,
            audio_input_size=audio_dim,
            visual_input_size=visual_dim,
            hidden_size=hidden_size,
            num_classes=num_classes
        ).to(DEVICE)

```

```

optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
criterion = nn.CrossEntropyLoss().to(DEVICE)

# Define checkpoint file name
checkpoint_file = os.path.join(
    checkpoint_dir, f"fold_{fold_num}_seed_{seed + 1}.pth")

# Check if checkpoint exists and load it
start_epoch = 0
if os.path.exists(checkpoint_file):
    try:
        checkpoint = torch.load(checkpoint_file, map_location=DEVICE)
        model.load_state_dict(checkpoint['model_state_dict'])
        optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
        start_epoch = checkpoint['epoch']
        last_loss = checkpoint.get('loss', 'N/A') # Get last saved
loss if available
        print(f"    Resuming training from checkpoint
{checkpoint_file} at epoch {start_epoch} (Last saved loss: {last_loss})")
    except Exception as e:
        print(f"    Warning: Could not load checkpoint
{checkpoint_file}: {e}. Starting from scratch.")
        start_epoch = 0

# Train and evaluate the model
epoch_losses = []
best_f1_seed = -1.0 # Track best F1 for this seed run

for epoch in range(start_epoch, num_epochs):
    accuracy, f1, loss = train_evaluate(
        model,
        nlp_train, audio_train, visual_train, labels_train,
        nlp_test, audio_test, visual_test, labels_test,
        optimizer, criterion, DEVICE, epoch)

    epoch_losses.append(loss)

# Optional: Save checkpoint only if F1 improves for this seed
if f1 > best_f1_seed:
    best_f1_seed = f1
    torch.save({
        'epoch': epoch + 1,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'loss': loss,
        'f1': f1, # Save F1 score in checkpoint
        'accuracy': accuracy,
    }, checkpoint_file)
    # print(f"    Epoch {epoch + 1}/{num_epochs} - Loss:
{loss:.4f}, Acc: {accuracy:.4f}, F1: {f1:.4f} (Checkpoint Saved)")
else:
    # print(f"    Epoch {epoch + 1}/{num_epochs} - Loss:
{loss:.4f}, Acc: {accuracy:.4f}, F1: {f1:.4f}")

# Store results from the *last* epoch for this seed
# (Alternatively, load the best checkpoint and evaluate on that)

```

```

        # For simplicity, using last epoch results here:
        fold accuracies_seeds.append(accuracy)
        fold_f1s_seeds.append(f1)
        fold_losses_seeds[seed + 1] = epoch_losses

    # Average metrics across seeds for this fold
    if fold accuracies_seeds: # Check if any seeds ran successfully
        avg_fold_accuracy = np.mean(fold accuracies_seeds)
        avg_fold_f1 = np.mean(fold_f1s_seeds)
        all accuracies.append(avg_fold_accuracy)
        all_f1s.append(avg_fold_f1)
        # Storing all epoch losses per seed per fold can be large - maybe
just avg loss?
        # For now, let's store the list of losses for the last epoch of each
seed
        last_epoch_losses = [losses[-1] for losses in
fold_losses_seeds.values() if losses]
        all_losses.append(np.mean(last_epoch_losses) if last_epoch_losses
else float('nan'))

    print(f" Fold {fold_num} Average (across seeds) - Accuracy:
{avg_fold_accuracy:.4f}, F1: {avg_fold_f1:.4f}")
    else:
        print(f" Fold {fold_num} - No successful seed runs.")

# Overall results across folds
if all accuracies:
    overall_avg_accuracy = np.mean(all accuracies)
    overall_avg_f1 = np.mean(all_f1s)
    print(f"\n--- Overall LOSO Results ---")
    print(f"Overall Average Accuracy: {overall_avg_accuracy:.4f}")
    print(f"Overall Average F1-score: {overall_avg_f1:.4f}")
    # print(f"Average last epoch losses across folds: {all_losses}")
else:
    print("\n--- No folds completed successfully. Cannot calculate overall
results. ---")

#--- 8. HierarchicalSpatioTemporalAttention ---
class HierarchicalSpatioTemporalAttention(nn.Module):
    """Hierarchical Spatio-Temporal Attention module for video processing."""
    def __init__(self, input_size, hidden_size, num_heads=8, dropout=0.1):
        super(HierarchicalSpatioTemporalAttention, self).__init__()

        # Multi-head attention for spatial features
        self.spatial_attention = nn.MultiheadAttention(
            embed_dim=input_size,
            num_heads=num_heads,
            dropout=dropout
        )

        # Temporal attention layers for different scales
        self.temporal_attention1 = nn.MultiheadAttention(
            embed_dim=input_size,
            num_heads=num_heads,
            dropout=dropout

```

```

    )
    self.temporal_attention2 = nn.MultiheadAttention(
        embed_dim=input_size,
        num_heads=num_heads,
        dropout=dropout
    )

    # LSTM for temporal encoding
    self.temporal_encoder = nn.LSTM(
        input_size=input_size,
        hidden_size=hidden_size,
        num_layers=2,
        batch_first=True,
        bidirectional=True,
        dropout=dropout
    )

    # Layer normalization
    self.layer_norm1 = nn.LayerNorm(input_size)
    self.layer_norm2 = nn.LayerNorm(input_size)
    self.layer_norm3 = nn.LayerNorm(input_size)

    # Dropout
    self.dropout = nn.Dropout(dropout)

    # Projection layers
    self.projection = nn.Linear(hidden_size * 2, input_size)

    def forward(self, x):
        # x shape: (batch_size, num_frames, num_features)
        batch_size, num_frames, num_features = x.size()

        # Reshape for spatial attention
        spatial_x = x.view(batch_size * num_frames, 1, num_features)
        spatial_x = spatial_x.transpose(0, 1) # (1, batch_size * num_frames,
num_features)

        # Apply spatial attention
        spatial_out, _ = self.spatial_attention(spatial_x, spatial_x, spatial_x)
        spatial_out = spatial_out.transpose(0, 1) # (batch_size * num_frames, 1,
num_features)
        spatial_out = spatial_out.view(batch_size, num_frames, num_features)
        spatial_out = self.layer_norm1(x + self.dropout(spatial_out))

        # First temporal scale (original frame rate)
        temporal_out1, _ = self.temporal_attention1(
            spatial_out.transpose(0, 1),
            spatial_out.transpose(0, 1),
            spatial_out.transpose(0, 1)
        )
        temporal_out1 = temporal_out1.transpose(0, 1)
        temporal_out1 = self.layer_norm2(spatial_out +
self.dropout(temporal_out1))

        # Second temporal scale (downsampled)
        if num_frames > 1:
            downsampled = temporal_out1[:, ::2, :] # Simple downsampling
            temporal_out2, _ = self.temporal_attention2(
                downsampled.transpose(0, 1),

```



```

        downsampled.transpose(0, 1),
        downsampled.transpose(0, 1)
    )
    temporal_out2 = temporal_out2.transpose(0, 1)
    temporal_out2 = self.layer_norm3(downsampled +
self.dropout(temporal_out2))

    # Upsample and combine
    temporal_out2 = F.interpolate(
        temporal_out2.transpose(1, 2),
        size=num_frames,
        mode='linear',
        align_corners=False
    ).transpose(1, 2)
    temporal_out = temporal_out1 + temporal_out2
else:
    temporal_out = temporal_out1

    # Final temporal encoding with LSTM
    lstm_out, _ = self.temporal_encoder(temporal_out)
    lstm_out = self.projection(lstm_out)

    return lstm_out

# --- 9. Main Execution Block ---
if __name__ == "__main__":
    start_main_time = time.time()

    # --- Configuration ---
    # !!! ADJUST THESE PATHS TO YOUR DATASET LOCATION !!!
    data_dir = 'Real-life_Deception_Detection_2016' # Example path
    annotation_file = "Real-
life_Deception_Detection_2016\Annotation\All_Gestures_Deceptive and Truthful.csv"
    checkpoint_dir = "multimodal_checkpoints"
    num_epochs_main = 50 # Adjust number of epochs
    learning_rate_main = 0.001
    hidden_size_main = 128 # Hidden dimension for feature processing/fusion

    # --- Workflow ---
    # 1. Load Data (Paths, Annotations, Transcriptions)
    data = load_data(data_dir, annotation_file)

    if not data:
        print("No data loaded. Exiting.")
        exit()

    # 2. Identify Subjects (Mandatory Facial Recognition)
    subject_mapping = identify_subjects_facial_recognition(data)
    # print("Subject Mapping (Facial Recognition):", subject_mapping) # Optional
print

    # 3. Prepare Data for LOSO (Get filtered annotations, raw transcriptions,
mapped IDs)
    annotations, transcriptions_raw, mapped_subject_ids, valid_indices =
prepare_loso(data, subject_mapping)

    if len(annotations) == 0:

```

```

    print("No valid trials remaining after preparing for LOSO. Exiting.")
    exit()

# Filter original data list based on valid indices from prepare_loso
# This ensures feature extraction only happens for trials included in LOSO
valid_data = [data[i] for i in valid_indices]
video_paths_valid = [item['video_path'] for item in valid_data]

# 4. Extract Features for the valid trials
nlp_features = extract_nlp_features(transcriptions_raw) # Takes raw text
audio_features = extract_audio_features(video_paths_valid)
visual_features = extract_visual_features(video_paths_valid)

# Sanity check feature shapes before running LOSO
print(f"Feature shapes: NLP={nlp_features.shape},
Audio={audio_features.shape}, Visual={visual_features.shape}")
if not (nlp_features.shape[0] == audio_features.shape[0] ==
visual_features.shape[0] == len(annotations)):
    print("Error: Feature array lengths do not match number of annotations
after filtering. Exiting.")
    print(f"Lengths: Annotations={len(annotations)},
NLP={nlp_features.shape[0]}, Audio={audio_features.shape[0]},
Visual={visual_features.shape[0]}")
    exit()

# 5. Run LOSO Cross-Validation with Multimodal Features
run_loso(annotations, nlp_features, audio_features, visual_features,
mapped_subject_ids,
        checkpoint_dir=checkpoint_dir,
        num_epochs=num_epochs_main,
        learning_rate=learning_rate_main,
        hidden_size=hidden_size_main)

end_main_time = time.time()
print(f"\nTotal execution time: {(end_main_time - start_main_time) / 60:.2f}
minutes")
print("Multimodal training complete.")

```

OUTPUT:

--- Overall LOSO Results ---

Overall Average Accuracy: 0.6249

Overall Average F1-score: 0.6200